

KRUTIK RAJESH PANCHAL

Rochester, NY | 585-957-2641 | krutikpanchal44@gmail.com | LinkedIn/panchalkrutik | Github/Krutik4

Education

| | |
|----------------------------------------------------------------|----------------------------|
| Rochester Institute of Technology (RIT) , Rochester, NY | Aug 2023 - May 2025 |
| <i>Master of Science in Data Science</i> | |
| University of Mumbai , Mumbai | Jul 2016 - May 2020 |
| <i>Bachelor of Engineering in Information Technology</i> | |

Skills

| |
|-------------------------------------------------------------------------------------------------------------------------------|
| Programming: Python (boto3, PySpark, scikit-learn, pandas, NumPy), SQL, Java |
| Databases: MySQL, PostgreSQL, MongoDB, Neo4j, Redshift |
| Cloud & DevOps: AWS, Azure, Docker, Kubernetes, Jenkins CI/CD , Git, GitHub, GitLab, Linux |
| Big Data: Apache Spark, Apache Airflow, Kafka |
| Visualization: Power BI(Microsoft Certified Data Analyst), Tableau, Qlik Sense, Excel, matplotlib, seaborn |
| Machine Learning & AI: Supervised/Unsupervised Learning, Feature Selection, Hyperparameter Tuning, NLP (LLMs, RAG) |

Experience

| | |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------|
| Data Engineer | Aug 2020 - Jul 2023 |
| <i>Larsen and Toubro Infotech</i> | <i>Mumbai, India</i> |
| <ul style="list-style-type: none">Architected and managed end-to-end data solutions on AWS using services like EC2, S3, Redshift, Glue, integrated with ETL tools like Talend, dbt resulting in a 40% improvement in data processing speedOptimized ETL and data warehouse processes by embedding automated validation checks, resulting in a 36% reduction in data discrepancies and improved data integrityImplemented real-time dashboards with Power BI and Tableau/ reducing processing time from 2 hours to 1 minute and accelerating data-driven decision-making across multiple projectsCollaborated with cross-functional teams and stakeholders in agile environments, transforming business needs into effective data-driven solutions that supported key project deliverables | |
| Graduate Teaching Assistant | Jan 2025 - May 2025 |
| <i>Rochester Institute of Technology</i> | <i>New York</i> |
| <ul style="list-style-type: none">Designed and validated ETL pipelines for batch and stream processing workflows using Python(PySpark), SQL, MongoDB, and Linux, reducing processing times by 30% and ensuring data integrityStreamlined deployment and automation by integrating GitLab CI/CD and Apache Airflow, accelerating deployment cycles by 25% while enhancing system scalability and error handlingMentored 33 students across 5 hands-on projects in Software Engineering for Data Science (Data Engineering), fostering expertise in scalable big data systems and end-to-end ETL pipelines | |

Projects

| | |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|
| LambdaSight : Serverless ETL & Analytics | Dec 2024 |
| <ul style="list-style-type: none">Created a serverless ETL pipeline using AWS Lambda, Glue, Python, and SQL to extract, aggregate, and load data from Amazon Redshift, cutting down execution time by 50%Leveraged AWS CloudWatch for monitoring, EventBridge for event-driven orchestration, and IAM for secure access control to automate data workflows, cutting manual tasks by 30%Developed interactive PowerBI dashboards that provided real-time insights for stakeholders, enhancing decision-making through optimized SQL queries and data transformations | |
| StockSage : Predictive Inventory Management | Jan 2025 |
| <ul style="list-style-type: none">Constructed a robust PySpark-based ETL solution to integrate data from MySQL, NoSQL(MongoDB Atlas, and JSON files) into medallion architecture, reducing processing time by 30% and ensuring high data accuracyOrchestrated efficient batch processing workflows with Apache Airflow to transform raw order transactions into structured tables for dynamic inventory management, leveraging Docker for containerized deploymentsIntegrated a time series forecasting model to predict sales and profit, delivering critical performance metrics (MAE, MSE) to support strategic decision-making | |
| DataNexus : Data Modelling & Dashboard | Mar 2025 |
| <ul style="list-style-type: none">Performed data transformation tasks via PySpark-driven workflows to process and normalize over 1M+ order records into a structured star-snowflake schema, improving data query efficiency and accessibilityApplied robust data validation, removing incomplete transactions and generating 10 core dimension and fact tables along with store-level partitioned datasets to support modular data analysis and reportingCreated an interactive Tableau dashboard showcasing KPIs like revenue trends, store-wise sales, supplier performance thus enhancing visibility for business stakeholders and reducing manual reporting time by 40% | |
| TickSync : Cloud-Native Stock Stream with Kafka & S3 | Apr 2025 |
| <ul style="list-style-type: none">Deployed an end-to-end real-time data pipeline using Apache Kafka on AWS EC2 to simulate stock trades from a 1M+ record dataset using a Python-based producer with Boto3Streamed data to Kafka consumer application which processes and stores incoming trade data to Amazon S3 in Parquet format, optimizing downstream processing and reducing query latency by 45%Automated metadata management with AWS Glue Crawlers and queried live trading metrics using Amazon Athena, achieving sub-second insights into price trends and trading volumes | |